

Assignment 1: Hash Tables

CS3D5A, Trinity College Dublin

Deadline: 17:30 12/10/2018

Grading: The assignment will be graded during the lab hours on 12/10/2018

Submission: Submit via blackboard. Include a separate .c file for each task, and the short assignment report in pdf, word, or text file, detailing your results for each part

Goals:

- Learn how to implement a hash table in C
- Consider how to choose a hash function
- Learn how to evaluate the performance of a hash table based on collisions
- Use a hash table to extract information from a dataset of historical figures

Task 1: Getting Started - 4 Marks

For the first part of this assignment you should write a hash table to **store the frequency** of names in an (unordered) list of Irish surnames. It will therefore use char arrays as keys and store ints as values.

Use the hash function hash1 that adds the integer value of the chars in the string:

```
int hash(char* s){
    int hash = 0;
    while(*s){
        hash = hash + *s;
        s++;
    }
    return hash;
}
```

You should use a linear probing strategy to handle collisions. For simplicity, you do not need to worry about a dynamically growing hash table. Allocate a fixed amount of memory for your keys and values at the start of the program and focus on implementing hashing and linear probing.

Test your hash table by loading the test data provided in names.csv. Write a program that allows you to test the frequency of a given string. NB: Load the data using a cvs parser (yours from assignment 0 or the one in the solution on blackboard) (if that doesn't work for you, keep making progress by hardcoding the data in your code!).

Update your code to count the number of collisions obtained and display it.

Listing 1: Sample output. Input from names.csv

```
names.csv loaded!
    Capacity   : 99991
    Num Terms  : 42
    Collisions : 0
    Load      : 0.000%
Enter term to get frequency or type "quit" to escape
>>> Synnot
Synnot 2
>>> bleb
bleb not in table
>>> quit
```

Task 1 Mark Allocations	
Correct implementation of hash table	3 marks
Take a string as input and print number of times string occurred in data	1 mark

Task 2: Choosing a hash function - 2 Marks

Now find a better hash function hash2 for the data considered. Feel free to consult online resources. Note, the sample data is only a sample! Do not overfit your function to the sample provided, it should work well with any lists of Irish surnames. Justify your choice of hash function in your report (half a page max). Test your function on the sample data, indicate in your report how many collisions occur, is it better than the result from task 1?

Task 2 Mark Allocations	
Justify your choice of hashing algorithm	1 mark
Implement and evaluate your hashing algorithm	1 mark

Task 3: Twice the Fun - 1 Mark

Choose a new hashing function hash3, and use it to augment your solution to Task 2 such that it uses double hashing instead of linear probing. As with the previous task, you should report the number of collisions which occurred in your hash table. Use this to demonstrate the improvement of double hashing over linear probing. Document in the report why you chose this function.

Listing 2: Sample output. Input from person.csv

```

.././data/people.csv loaded!
  Capacity   : 99991
  Num Terms  : 86077
  Collisions : 179986
  Load       : 0.861%
Enter term to get frequency or type "quit" to escape
>>> Maguire
Maguire 104
>>> Magwire
Magwire 141
>>> quit

```

Task 3 Mark Allocations	
Successful implementation of second hashing function	1 mark

Task 4: A More Interesting Application - 3 Marks

Hash tables have numerous applications in computer science. One domain where they can find great use is that of Information Retrieval. Search terms extracted from a collection of documents are used as the keys and lists of documents are stored as values in the table. When a user issues a query, the list of all documents which may be of interest to them can be rapidly retrieved, ranked and presented.

For this task you have been given a file containing a list of people. The data which you have been given is real data produced with great effort and expense by a number of Trinity historians. The people mentioned are all individuals who were in some way involved with the 1641 Irish rebellion. Learning about these people is of great interest to historians, but the challenging nature of working with 17th century data can make this difficult.

For this task you should expand (a copy of) your solution to the previous tasks. Instead of storing term counts at each index, you should store a list of people with a given surname. The keys of the hash table will be surnames. The values will be linked lists of people.

As before you should provide a way to search for information in the hash table. Allow a user to enter a surname and get a list of people with the given surname as a search result. Test on the truncated data and then on the full dataset.

Listing 3: Sample output.

```
.././data/people.csv loaded!
Capacity   : 99991
Num Terms  : 14963
Collisions : 2724
Load       : 0.150%
Enter term to get frequency or type "quit" to escape
>>> Poulton
Person ID Deposition ID      Surname      Forename Age Person Type
      3700      818191r164      Poulton      Anthony  0   Mentioned
      3678      818189r163      Poulton      Anthony  0   Mentioned
      3664      818187r162      Poulton      Anthony  0   Mentioned
      3576      818185r161      Poulton      Anthony  0   Deponent
      3428      818161r152      Poulton      Anthony  0   Mentioned
       716      819112r142      Poulton      Anthony  0   Mentioned
>>> Lawles
Person ID Deposition ID      Surname      Forename Age Person Type
      49837      823038r036      Lawles      Phillip  0      Debtor
      45499      812318r259      Lawles      Elizabeth 0      Deponent
>>> quit
```

Task 4 Mark Allocations	
Alter (a copy of) your previous solution so that it now stores lists of people as values instead of word counts. This alteration should include freeing any dynamically allocated memory required for the lists.	2 marks
Provide a means to search across your hash table of people	1 mark

Extra Bits - Not Graded

- As previously mentioned, the people in this dataset are real people. Do a search for your family name and see if anything comes up. If you find someone you would like to read about, the text of the depositions is available at <http://cultura-project.eu/1641>. You will need to register with the site. You can then search for the depositions using the deposition ID.
- You may have noticed that some people's names are spelled "wrong", e.g. Maguire spelled "Magwire". The English language was not standardised until some time in the mid-18th century. What you are seeing is the early Anglicization of Irish names coupled with the fact that the rules of the English language were not yet well established. Can you think of a way to provide more tolerant search across the hash table so that people searching for "Maguire" might also get results for "Magwire"?